MIDS-W261-EndOfTermExam-Published

December 7, 2016

1 W261 MIDS Machine Learning at Scale

1.1 End of Term exam

Week 14
Fall, 2016
Student Name: [Anthony, Spalvieri-Kruse]
Student ID: 26990271
Email: ask@ischool.berkeley.edu

2 Exam Schedule (All times are in California Time)

4:00 PM - 6:00 PM

3 Instructions for exam

1. Please acknowledge receipt of exam by sending a quick reply to the me
2. Review the submission form first to scope it out (it will take a 5-10 minutes to input your answers and other information into this form): http://goo.gl/forms/ggNYfRXz0t
3. Please keep all your work and responses in ONE (1) notebook only (and submit via the submission form)
4. Please make sure that the NBViewer link for your Submission notebook works
5. Please do NOT discuss this exam with anyone (including your class mates) until after Monday, Midday, of week 16
6. Please submit your solutions and notebook via the following form: http://goo.gl/forms/ggNYfRXz0t

This is an open book exam meaning you can consult webpages and textbooks, class notes, slides etc. but you can not each other or any other person/group. Please complete this exam by yourself within the time limit.

7. For markdown help in iPython Notebooks please see: https://sourceforge.net/p/ipython/discussion/markdown_syntax

```
In [1]: import os
        import sys

        #Change SPARK_HOME to point to the folder where you installed Spark
        spark_home = os.environ['SPARK_HOME'] = '/usr/local/spark'

        if not spark_home:
```

```
            raise ValueError('SPARK_HOME enviroment variable is not set')
        sys.path.insert(0,os.path.join(spark_home,'python'))
        sys.path.insert(0,os.path.join(spark_home,'python/lib/py4j-0.10.3-src.zip'))
        execfile(os.path.join(spark_home,'python/pyspark/shell.py'))

Welcome to
      ____              __
     / __/__  ___ _____/ /__
    _\ \/ _ \/ _ `/ __/  '_/
   /__ / .__/\_,_/_/ /_/\_\   version 2.0.2
      /_/

Using Python version 2.7.11 (default, Dec  6 2015 18:57:58)
SparkSession available as 'spark'.
```

4 Exam questions begins here

4.1 ET:1

Assume your tasked with modeling a REGRESSION problem. How do you determine which variables may be important?

Select the single most correct response from the following:

- (a) 1
- (b) 2
- (c) 3
- (d) 1, 2, 3, 4

Answer: d

4.2 ET:2

Using one-hot-encoding, a categorical feature with four distinct values would be represented by how many features?

(a) 1 feature

(b) 2 features

(c) 4 features

(d) none of the above

Answer: c) each value is a feature

4.3 ET:3

In the following (and also referring to HW12: http://nbviewer.jupyter.org/urls/dl.dropbox.com/s/1wb2rdqbet54y1h/MIDS-MLS-Project-Criteo-CTR.ipynb) we have hashed the three sample points using numBuckets=4 and numBuckets=100. Complete the three statements below about these hashed features summarized in the following table using each answer once.

Name	Raw Features	4 Buckets	100 Buckets
sampleOne	[(0, 'mouse'), (1, 'black')]	{2: 1.0, 3: 1.0}	{14: 1.0, 31: 1.0}
sampleTwo	[(0, 'cat'), (1, 'tabby'), (2, 'mouse')]	{0: 2.0, 2: 1.0}	{40: 1.0, 16: 1.0, 62: 1.0}
sampleThree	[(0, 'bear'), (1, 'black'), (2, 'salmon')]	{0: 1.0, 1: 1.0, 2: 1.0}	{72: 1.0, 5: 1.0, 14: 1.0}

With 100 buckets, sampleOne and sampleThree both contain index 14 due to _____.

(a) A hash collision

(b) Underlying properties of the data

(c) The fact that used 100 buckets

(d) none of the above

Answer: b; they share a raw feature

4.4 ET:4

In the following (and also referring to HW12: http://nbviewer.jupyter.org/urls/dl.dropbox.com/s/1wb2rdqbet54y1h/MIDS-MLS-Project-Criteo-CTR.ipynb) we have hashed the three sample points using numBuckets=4 and numBuckets=100. Complete the three statements below about these hashed features summarized in the following table using each answer once.

Name	Raw Features	4 Buckets	100 Buckets
sampleOne	[(0, 'mouse'), (1, 'black')]	{2: 1.0, 3: 1.0}	{14: 1.0, 31: 1.0}

Name	Raw Features	4 Buckets	100 Buckets
sampleTwo	[(0, 'cat'), (1, 'tabby'), (2, 'mouse')]	{0: 2.0, 2: 1.0}	{40: 1.0, 16: 1.0, 62: 1.0}
sampleThree	[(0, 'bear'), (1, 'black'), (2, 'salmon')]	{0: 1.0, 1: 1.0, 2: 1.0}	{72: 1.0, 5: 1.0, 14: 1.0}

It is likely that sampleTwo has two indices with 4 buckets, but three indices with 100 buckets due to _____.

(a) A hash collision

(b) Underlying properties of the data

(c) The fact that we go from 4 to 100 buckets

(d) none of the above

Answer: c

4.5 ET:5

In the following (and also referring to HW12: http://nbviewer.jupyter.org/urls/dl.dropbox.com/s/1wb2rdqbet54y1h/MIDS-MLS-Project-Criteo-CTR.ipynb) we have hashed the three sample points using numBuckets=4 and numBuckets=100. Complete the three statements below about these hashed features summarized in the following table using each answer once.

Name	Raw Features	4 Buckets	100 Buckets
sampleOne	[(0, 'mouse'), (1, 'black')]	{2: 1.0, 3: 1.0}	{14: 1.0, 31: 1.0}

Name	Raw Features	4 Buckets	100 Buckets
sampleTwo	[(0, 'cat'), (1, 'tabby'), (2, 'mouse')]	{0: 2.0, 2: 1.0}	{40: 1.0, 16: 1.0, 62: 1.0}
sampleThree	[(0, 'bear'), (1, 'black'), (2, 'salmon')]	{0: 1.0, 1: 1.0, 2: 1.0}	{72: 1.0, 5: 1.0, 14: 1.0}

With 4 buckets, sampleTwo and sampleThree both contain index 0 due to _____.

(a) A hash collision

(b) Underlying properties of the data

(c) The fact that we use 4 buckets

(d) none of the above

Answer: a

4.6 ET:6 When applying numerical machine learning approaches (and for non-numerical approaches if required) to big data problems which of the following steps are could be used during modeling and are recommended:

(a) Convert categorical features to numerical features via one-hot-encoding and store in a dense representation

(b) Transform categorical features using hashing regardless of how many unique categorical values exist in training and test data

(c) Use matrix factorization to remap your input vectors to latent concepts

(d) none of the above

Answer: a, b, c

4.7 ET:7

When dealing with numercial data which of the following are ways to deal with missing data:

(a) Delete records that have missing input values

(b) Standardize the data and set all missing values to 1 (one)

(c) Use K-nearest neighbours based on the test set to fill in missing values in the training set

(d) none of the above

Answer: Technically a, b, and c are approaches that could have tangible benefits, but they're all flawed. With a) we're deleting data that could have utility, limiting our insights. In b) arbitrarily setting an input to 1 may have a huge outlying effect on the model. With c) we have to go through the trouble of creating a kNN model at scale, along with filling in the training set with test set data which is a terrible idea.

4.8 ET:8

In the Criteo project, we're trying to predict what:

(a) Revenue from click events

(b) Click-through vs not click event

(c) Probability of a purchase

(d) none of the above

Answer: b

4.9 ET:9

Which of the following are true about the purpose of a loss function?

(a) It's a way to penalize a model for incorrect predictions

(b) It precisely defines the optimization problem to be solved for a particular learning model

(c) Loss functions can be used for modeling both classification and regression problems

(d) none of the above

Answer: a, b, c

4.10 ET:10

Many standard machine learning methods can be formulated as a convex optimization problem, i.e. the task of finding a minimizer of a convex function f that depends on a variable vector **w** (called weights in the code), which has d entries. Formally, we can write this as the optimization problem minw∈ℝdf(w), where the objective function is of the form (JGS)

Here the vectors xi∈ℝd are the training data examples, for 1≤i≤n, and yi∈ℝ are their corresponding labels, which we want to predict. We call the method linear if L(w;x,y) can be expressed as a function of _ __wT__x _ and y. Several of spark.mllib's classification and regression algorithms fall into this category.

The objective function **f** has two parts: the regularizer that controls the complexity of the model, and the loss that measures the error of the model on the training data. The loss function L(w;.) is typically a convex function in **w**. The fixed regularization parameter $\lambda \geq 0 \lambda \geq 0$ (regParam in the code) defines the trade-off between the two goals of minimizing the loss (i.e., training error) and minimizing model complexity (i.e., to avoid overfitting).

When implementing Logistic (or linear) Regression with Regularization in Spark which of the following apply when using the above cost functions (mulitple options may apply):

(I) When lambda equals one, it provides the same result as standard logistic/linear regression

(II) One only needs to modify the standard logistic (linear) regression (i.e., with no regularization term) by adding some code after the map-reduce (loss) gradient steps

(III) When lambda equals zero, it provides the same result as standard logistic (linear) regression
(IV) None of the above

Select the most correct from the following:

- (a) I, II, III
- (b) II, III
- (c) III
- (d) IV

Answer: b ; a is definitely wrong, b is def right because I was embarrased in class about this and have not forgotten, and c should be apparent because then the R(x) term vanishes.

4.11 ET:11

In the context of ecommerce you have just deployed a new conversion rate prediction model to production. This model (aka treatment model) will challenge the control nodel (i.e., the current model) in AB Test manner to see if it can be produce better revenue. Here is the data that was taken from this live AB Test.

```
CONTROL MODEL (our new CTR model)
Impression ID   Revenue
1                $0.50
2                $0.50
3                $3.00
......
20000            $3.00
20001            $3.00
20002            $3.00
20003            $3.00
......
50,001           $3.00
.....
100,000          $4.00
```

All other impressions in this 100,000 sample resulted in zero transactions and therefore zero revenue.

```
TREATMENT MODEL (our new CTR model)
Impression ID   Revenue
1                $1.50
2                $0.50
3                $0.00
......
50,001           $3.00
.....
100,000          $4.00
```

All other impressions in this 100,000 sample resulted in zero transactions and therefore zero revenue.

P-values are a common way to determine the statistical significance of a test. The smaller it is, the more confident you can be that the test results are due to something other than random chance. A common

p-value of .05 is a 5% significance level. Similarly, a p-value of .01 is a 1% significance level. A p-value of .20 is a 20% significance level. For this problem set the p-value to 0.01

Which of the following are true:

(a) Based on revenue there is no statistical significant difference between the Control and the Treatment at p-value of 0.05 for a one-sided t-test

(b) Based on transaction rates (tranactions that generated revenue versus not) there is no statistical significant difference between the Control and the Treatment at p-value of 0.05 for a one-sided t-test

(c) AB testing using differences in revenue for this problem is a useful means of determining if the Treatment conversion rate prediction model is better than the control model.

(d) none of the above

```
In [1]: import numpy as np
        from scipy.sparse import coo_matrix
        import scipy.stats

        control=[.5, .5, 3, 3, 3, 3, 3, 3, 4]
        control.extend([0]*(100000-len(control)))
        control=np.asarray(control)

        treatment=[ 0, 0, 0, 0, 0, 0, 0, 0, 0]
        treatment.extend([0]*(100000-len(treatment)))
        treatment=np.asarray(treatment)

        #nb 2-sided test
        scipy.stats.ttest_ind(control, treatment), scipy.stats.ttest_ind(control.astype(bool), treatmen

Out[1]: (Ttest_indResult(statistic=2.7393492169656311, pvalue=0.0061566398814161729),
         Ttest_indResult(statistic=3.000120008400637, pvalue=0.0026990648570512994))
```

Answer: c; So, the multiple choice lists a p value of .05 though the body says .01. It shouldn't matter either way though. We're doing a 2 sided test, so we want to see that $p/2 < .01$. Clearly it is in each test, meaning that we have statistically significant differences between the two sample means for both revenue and raw transaction. Only c is true.

4.12 ET:12

Given this graph expressed in the form of an adjacency list,

```
Node  adjacentNode:weightAssociatedWithEdge
N1    N6:10,  N2:2
N2    N3:1
N3    N4:1
N4    N5:1
N5    N6:1
N6    N7:1
N7    N8:1
N8    N9:1
```

Using the parallel breadth-first search algorithm for determining the shortest path from a single source, how many iterations are required to discover the shortest distances to all nodes from Node 1

A 7

B 8

C 13
D None of the above

Answer: b; Know we have to traverse whole graph because 2 + 7 traversals is distance of 9 to get to 6, which is less than first path taking 10 units.

```
n1 -> {N6, N2} | {N6: 10, N2: 2} | {N6, N2} on deck
n6 -> {N7}     | {N6: 10, N2: 2, N7: 11} | {N2, N7} on deck
n2 -> {N3}     | {N6: 10, N2: 2, N7: 11, N3:3} | {3, 7} OD
n7 -> {n8}     | ... | {8, 3} OD
n3 -> {n4}     | ... | {8, 4} OD
n8 -> {n9}     |     | {4} OD
n4 -> {n5}     |     | {5} OD
n5 -> {n6}     |     | {} OD --> end at 8 rounds
```

4.13 ET:13 Assume the Lagrangian for SOFT SVMs (unconstrained optimization) is as follows:

minimize [λ/2 * w'w + CΣi=1:n ξi + 1/n Σi=:n max(0, 1 - ξi – yi(w'xi – b)))]

(a) When λ is super small (e.g., 0.000001), then the above Langrangian will yield a Hard SVM

(b) In the context of support vector machines, linear kernels can be readily parallelized in map reduce frameworks such as Spark

(c) Sequential learning via algorithms such perceptron can take advantage of map-reduce frameworks and yield exactly the same results as a single core implementation with significant reductions in training time

(d) When λ is 1.0, then the above Langrangian will yield a Soft SVM

Answer: a,b,c,d

4.14 ET:14 Given the following paired RDDs

RDD1 = {(1, 2), (3, 4), (3, 6)} RDD2 = {(3, 9) (3, 6)}

Using PySpark, write code to perform an inner join of these paired RDDs. What is the resulting RDD? Make your Spark available in your notebook:

A: [(3, (4, 9)), (3, (6, 9))]
B: [(3, (4, 9)), (3, (4, 6)), (3, (6, 9)), (3, (6, 6))]
C: [(3, (4, 9)), (3, (4, 6)), (3, (6, 9)), (3, (6, 9))]
D: None of the above

```
In [5]: rdd1 = spark.sparkContext.parallelize({(1, 2), (3, 4), (3, 6)})
        rdd2 = spark.sparkContext.parallelize({(3, 9),(3, 6)})
        rdd1.join(rdd2).collect()

Out[5]: [(3, (4, 9)), (3, (4, 6)), (3, (6, 9)), (3, (6, 6))]
```

Answer: B

4.15 ET:15 You have been tasked to build a predictive model to forecast beer sales for a chain of stores.

After doing basic exploratory analysis on the data, what is the first thing you do regarding modeling?

(a) Construct a baseline model

(b) Determine a metric to evaluate your machine learnt models

(c) Split your data into training, validation and test subsets (or split using cross fold validatation)

(d) All of the of the above

Answer: d

4.16 ET:16

Use Spark and the following notebook to answer this question:

- http://nbviewer.jupyter.org/urls/dl.dropbox.com/s/6s5ph41h74bggwi/Linear-Regression-on-Beer-Data.ipynb
- https://www.dropbox.com/s/6s5ph41h74bggwi/Linear-Regression-on-Beer-Data.ipynb?dl=0

The mean absolute percentage error (MAPE), also known as mean absolute percentage deviation (MAPD), is a measure of prediction accuracy of a model for say a forecasting method in statistics, for example in trend estimation. It usually expresses accuracy as a percentage, and is defined by the formula:

MAPE = average over all examples (100*Abs(Actual - Predicted) / Actual))

Note when Actual is zero that test row is dropped from the evaluation.

Construct a mean model for target variable `CASES18PK`. Calculate the MAPE for the mean model over the training set. Select the closest answer.

(a) 200%

(b) 250%

(c) 20%

(d) 180%

```
In [6]: %%writefile beerSales.txt
        Week        PRICE12PK        PRICE18PK        PRICE30PK        CASES12PK        CASES18PK
        1        19.98        14.10        15.19        223.5        439        55.00
        2        19.98        18.65        15.19        215.0        98        66.75
        3        19.98        18.65        13.87        227.5        70        242.00
        4        19.98        18.65        12.83        244.5        52        488.50
        5        19.98        18.65        13.16        313.5        64        308.75
        6        19.98        18.65        15.19        279.0        72        111.75
        7        19.98        18.65        13.92        238.0        47        252.50
        8        20.10        18.73        14.42        315.5        85        221.25
        9        20.12        18.75        13.83        217.0        59        245.25
        10        20.13        18.75        14.50        209.5        63        148.50
        11        20.14        18.75        13.87        227.0        57        229.75
        12        20.12        18.75        13.64        216.5        54        312.00
        13        20.12        13.87        14.31        169.0        404        96.75
        14        20.13        14.27        13.85        178.0        380        123.25
        15        20.14        18.76        14.20        301.5        65        200.50
        16        20.14        18.77        13.64        266.5        40        359.75
        17        20.13        13.87        14.33        182.5        456        113.50
        18        20.13        14.14        13.14        159.0        176        136.50
        19        20.13        18.76        13.81        285.5        61        225.50
        20        20.13        18.72        15.19        360.0        91        122.25
        21        20.13        18.76        13.13        263.0        59        443.75
```

```
22        19.18        18.76        13.63        443.5        83        322.75
23        14.78        18.74        15.19        1101.5        41        53.00
24        16.04        18.75        13.89        814.0        47        140.75
25        20.12        18.75        14.28        365.0        84        210.75
26        19.75        18.75        15.19        510.0        85        110.50
27        19.65        18.75        13.12        580.5        116        568.25
28        19.69        13.79        13.78        251.0        544        115.50
29        20.12        13.49        15.19        237.0        890        58.75
30        20.12        14.89        15.19        302.5        371        77.25
31        20.13        13.94        15.19        229.5        557        66.25
32        20.14        13.67        15.19        188.5        775        50.00
33        15.14        14.43        15.19        795.5        236        46.50
34        14.33        18.75        15.19        1556.5        43        65.75
35        16.24        18.22        13.14        807.5        63        252.75
36        19.93        14.06        13.45        243.0        469        179.00
37        21.06        14.43        13.00        201.5        335        226.25
38        21.19        19.48        13.60        294.0        75        288.50
39        21.23        15.15        14.46        220.5        461        114.25
40        20.12        13.79        14.94        255.5        817        70.00
41        14.73        14.31        15.19        920.5        200        47.75
42        14.57        19.50        15.19        730.0        32        98.75
43        15.94        13.85        15.19        262.5        460        77.00
44        20.70        14.23        13.43        209.5        751        160.50
45        19.57        19.31        14.37        283.0        70        143.50
46        19.60        19.29        15.19        262.5        80        133.00
47        19.94        13.76        15.19        310.0        523        68.75
48        21.28        13.45        15.19        278.5        741        81.75
49        14.56        15.13        15.19        741.5        130        56.25
50        14.39        19.43        15.19        1316.0        69        68.75
51        16.81        13.26        15.19        449.0        493        49.25
52        19.86        13.92        15.19        505.0        814        76.50

Writing beerSales.txt
```

```
In [29]: df=sc.textFile("beerSales.txt")

         header = df.first() #extract header
         df = df.filter(lambda x: x != header)

         meanParts = df.map(lambda x: (1,float(x.split('\t')[5]))).reduce(lambda x, y: [x[0]+y[0], x[1]+

         mean = spark.sparkContext.broadcast(meanParts[1]/meanParts[0])

         mape = df.map(lambda x: float(x.split('\t')[5])).map(lambda x: 100*abs(x-mean.value)/x).reduce

         print mape
```

```
200.44422333
```

Answer: a; I'm interpreting 'mean model' to mean setting every prediction equal to the mean of the sample.

4.17 ET:17

Use Spark and the following notebook to answer this question:

- http://nbviewer.jupyter.org/urls/dl.dropbox.com/s/6s5ph41h74bggwi/Linear-Regression-on-Beer-Data.ipynb
- https://www.dropbox.com/s/6s5ph41h74bggwi/Linear-Regression-on-Beer-Data.ipynb?dl=0

The target variable `CASES18PK` is skewed, so take the log of it (and make it more normally distributed) and compute the MAPE of the mean model for `CASES18PK`. Select the closest answer to your calculated MAPE.

(a) 200%

(b) 30%

(c) 20%

(d) 10%

```
In [30]: import math

         df=sc.textFile("beerSales.txt")
         df = df.filter(lambda x: x != header)

         meanParts = df.map(lambda x: (1,float(x.split('\t')[5]))).reduce(lambda x, y: [x[0]+y[0], x[1]+
         meanLogParts = df.map(lambda x: (1,math.log(float(x.split('\t')[5])))).reduce(lambda x, y: [x[

         mean = spark.sparkContext.broadcast(meanParts[1]/meanParts[0])
         logMean = spark.sparkContext.broadcast(meanLogParts[1]/meanLogParts[0])

         mapeLogTargetOnly = df.map(lambda x: math.log(float(x.split('\t')[5]))).map(lambda x: 100*abs(
         mapeLogTargetAndLogMean = df.map(lambda x: math.log(float(x.split('\t')[5]))).map(lambda x: 10


         print "mapeLogTargetOnly: ", mapeLogTargetOnly
         print "mapeLogTargetAndLogMean: ", mapeLogTargetAndLogMean
```

```
mapeLogTargetOnly:  5237.79963292
mapeLogTargetAndLogMean:  19.5849342896
```

Answer: c; I'm assuming that you wanted to also recompute the mean model so that it's the mean of log(target)

4.18 ET:18

Use Spark and the following notebook to answer this question:

- http://nbviewer.jupyter.org/urls/dl.dropbox.com/s/6s5ph41h74bggwi/Linear-Regression-on-Beer-Data.ipynb
- https://www.dropbox.com/s/6s5ph41h74bggwi/Linear-Regression-on-Beer-Data.ipynb?dl=0

Build a linear regression model using the following variables:
Log(CASES18PK) ~ log(PRICE12PK), log(PRICE18PK), log(PRICE30PK)
Calculate MAPE over the test set and select the closest answer.

(a) 4.3%

(b) 4.6%

(c) 3.5%

(d) 3.9%

```
In [50]: from pyspark.mllib.regression import LabeledPoint, LinearRegressionWithSGD, LinearRegressionMo

         df=sc.textFile("beerSales.txt")

         df.take(10)
         df = df.filter(lambda x: x != header).map(lambda x: [float(i) for i in x.split('\t')])

         data = df.map(lambda x: LabeledPoint(math.log(x[5]),[math.log(x[1]),math.log(x[2]),math.log(x[
         reg=LinearRegressionWithSGD.train(data, iterations=10000, step=.01)
         preds=data.map(lambda x: (x.label, reg.predict(x.features)))
         mape = preds.map(lambda (v, p): 100*abs(v - p)/v)\
                     .reduce(lambda x, y: x + y)/preds.count()

         print mape
```

```
19.8406686036
52
```

b

4.19 ET:19

Recall that Spark automatically sends all variables referenced in your closures to the worker nodes. While this is convenient, it can also be inefficient because (1) the default task launching mechanism is optimized for small task sizes, and (2) you might, in fact, use the same variable in multiple parallel operations, but Spark will send it separately for each operation. As an example, say that we wanted to write a Spark program that looks up countries by their call signs (e.g., the call sign for Ireland is EJZ) by prefix matching in an table. In the following the "signPrefixes" variable is essentially a table with two columns "Sign" and "Country Name". The goal is to join the following tables:

`signPrefixes` table with columns "Sign" and "Country Name"
`contactCounts` table with columns "Sign" and "count"
to yield a new table:
`countryContactCounts` with the following columns "Country Name" and "count"
Use Spark and the following notebook to answer this question:

- http://nbviewer.jupyter.org/urls/dl.dropbox.com/s/6s5ph41h74bggwi/Linear-Regression-on-Beer-Data.ipynb
- https://www.dropbox.com/s/6s5ph41h74bggwi/Linear-Regression-on-Beer-Data.ipynb?dl=0

How can we modfify this code to make it more efficient? Choose one response only

(a) modify line 18 with `sc.broadcast(loadCallSignTable())`

(b) Use accumulators to store the counts for each country

(c) The code is already optimal

(d) none of the above

Answer:
a

```
In [ ]:
```